

cse5441 - parallel computing

https://cdn8.bigcommerce.com/s-o1bh36yyl/images/stencil/1280x1280/products/18124/29821/hpgraphiccardsvideocardtesla__52934.1516995900.jpg?c=2&imbypass=on
<https://i0.wp.com/www.visitcocosisland.com/wp-content/uploads/2017/01/animalia-life-club.jpg?resize=640%2C492>
<https://qph.fs.quoracdn.net/main-qimg-257ea3b2d296c1605a50e5fd5db21cfa-c>
<https://pictures.topspeed.com/IMG/jpg/201205/plymouth-hemi-cuda-p.jpg>



'cuda




why cuda?

- robust development environment
- robust developer community
- significant market share
- leading feature development
- used by Ohio Supercomputer Center
- used by OSU High Performance Computing group
- quickly evolving environment

Nvidia developer's forum: <https://devtalk.nvidia.com/>

Nvidia CUDA toolkit: <http://docs.nvidia.com/cuda/index.html>

10-Year Total Returns: NVIDIA vs. Select Competitors/Peers & the S&P 500



Desktop PC and AIBs (M units)





Figure 3: Shipments of desktop PCs compared to desktop AIBs over time

Discrete Desktop GPU Market Share


Data by Jon Peddie Research



what is a gpu?



"a single chip processor with integrated transform, lighting, triangle setup/clipping, and rendering engines that is capable of processing a minimum of 10 million polygons per second."



applications:

Abacus/Standard
Ablinit
AcuSolve
ADF
Amber
ArrayFire
AxRecon
COSMO
Desmond
Empro
Fluent
Optistruct
RADIOSS

Savant
Semcad-X
TeraChem
Theano
Torch7
Tsunami RTM
UGENE
Vega ZZ
Wolfram



application fields:


computational finance
computational physics
database/data science
design automation
fluid dynamics
media & entertainment
medical imaging
molecular dynamics
numerical analysis
quantum chemistry
structural mechanics
weather & climate

what is a gpu?



gpu / cpu comparison

"CPU"



GPU



computational

memory-related

pthread -vs- OpenMP -vs- cuda



pthread

- rigorous analysis of variable scope and sharing
- host processing
- global shared memory
- asynchronous peer threads


OpenMP

- rigorous analysis of variable scope and sharing
- host processing
- global shared memory
- asynchronous peer threads


- rigorous analysis of variable scope and sharing
- device / host processing
- non-contiguous, non-homogenous memory
- hierarchical, synchronous thread teams

cpu global shared memory

review



cpu / gpu non-contiguous memory



— “memory speed”

— “parallel-access memory speed”

cuda terminology

- **host** *a general-purpose computer (CPU) with which a GPU is associated*
- **device** *a GPU which is associated with a CPU*
- **(cuda) program** *a program written with both host and GPU (kernel) portions*
- **kernel** *the GPU portion of a (cuda) program*
- **barrier** *various synchronization methods which control the progress of threads*
- **thread hierarchy** *an organizational method by which threads are organized*
- **grid** and **block** *dimensional aspects of a thread hierarchy*

cuda terminology

- **thread** *unit of kernel execution with a GPU*
- **thread index** *the coordinates of a specific thread within a thread hierarchy*
- **thread id** *the linearized value of the thread index*
- **dim3** *a structured type used to store the grid and block sizes, each of which themselves may be in 1, 2 or 3 dimensions*
- **SM** *“streaming multiprocessor,” low clock, tightly coupled, high data-width processor group*
- **warp** *unit of kernel scheduling, a group of threads which are executed simultaneously on a specific SM*

cuda kernel threads

a kernel is a data-parallel function

- invoking a kernel creates **super-lightweight threads**
- each thread has a **context** within the kernel
- threads are generated and scheduled **by hardware**
- Nvidia features **0-cycle thread swap**

basic cuda application structure

ON HOST:

- instantiate host data structures
- perform initialization and pre-process data
- instantiate device data structures
- initialize device data structures
- define device thread hierarchy
- launch kernel
-
- copy results from device to host data structures
- de-allocate device data structures
-
- perform single-threaded application functions
-
- instantiate device data structures
- initialize device data structures
- define device thread hierarchy
- launch kernel
-
- copy results from device to host data structures
- de-allocate device data structures
-
- perform single-threaded application functions

(wash, rinse, repeat ...)


ON DEVICE:

cuda memory transfers

device side

device threads can:

- R/W per-thread registers
- R/W per-thread local memory (not shown)
- R/W per-block shared memory
- R/W per-grid “global” memory
- R-only per-grid “constant” memory



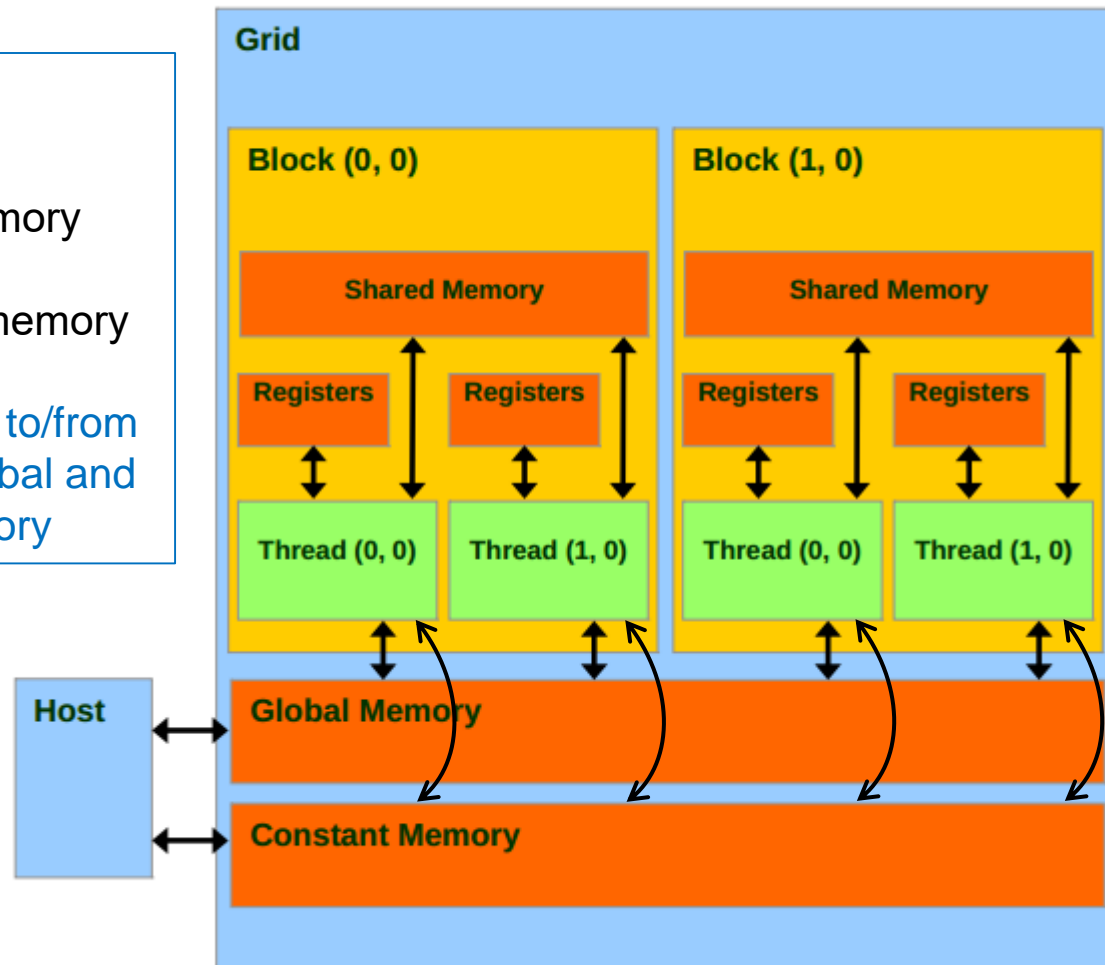
cuda memory transfers

host side

Host functions can:

- R/W per-grid “global” memory
- R/W per-grid “constant” memory

transfer memory contents to/from
host memory to/from global and
constant GPU memory



cuda memory transfers - allocation

host program

```
int *h_a;           //declare pointer for host memory
int *d_a;           //declare pointer for device memory
```

```
// allocate host and device memory
size_t memSize;
```

sizeof "whatever"
int as an example
v

```
memSize = num_blocks * num_th_per_blk * sizeof(int);
h_a      = (int*) malloc(memSize);
cudaMalloc( (void**) &d_a, memSize);
```

cuda memory transfers - copy

host program

```
int *h_a;           //declare pointer for host memory
int *d_a;           //declare pointer for device memory

// allocate host and device memory
size_t memSize;

memSize  = num_blocks * num_th_per_blk * sizeof(int);
h_a      = (int*) malloc(memSize);
cudaMalloc( (void**) &d_a, memSize);

// first put interesting stuff into h_a, then ...
cudaMemcpy( d_a, h_a, memSize, cudaMemcpyHostToDevice);
```

cuda memory transfers - retrieve

host program

```
int *h_a;           //declare pointer for host memory
int *d_a;           //declare pointer for device memory

// allocate host and device memory
size_t memSize;

memSize  = num_blocks * num_th_per_blk * sizeof(int);
h_a      = (int*) malloc(memSize);
cudaMalloc( (void**) &d_a, memSize);

cudaMemcpy( d_a, h_a, memSize, cudaMemcpyHostToDevice);

(launch kernel)

cudaMemcpy( h_a, d_a, memSize, cudaMemcpyDeviceToHost);
cudaFree(d_a);
```


cuda memory transfers

host program

```
int *d1_a;           //declare pointer for device memory
int *d2_a;           //declare another pointer for device memory
```

```
cudaMemcpy( d2_a, d1_a, memSize, cudaMemcpyDeviceToDevice);
```

```
int *h1_a;           //declare pointer for host memory
int *h2_a;           //declare another pointer for host memory
```

```
cudaMemcpy( h2_a, h1_a, memSize, cudaMemcpyHostToHost);
```

launching kernel

```
// launch kernel "initArray"
```

```
dim3 dimGrid(num_blocks);  
dim3 dimBlock(num_th_per_blk);  
initArray<<< dimGrid, dimBlock >>>(d_a);
```

a simple cuda program

// initialize an array using GPU

```
int main()
{
    int *h_a;           //pointer for host memory
    int *d_a;           //pointer for device memory
```

// define thread hierarchy

```
int num_blocks      = 8;
int num_th_per_blk  = 8;
```

// allocate host and device memory

```
size_t memSize;
```

```
memSize = num_blocks * num_th_per_blk * sizeof(int);
h_a      = (int*) malloc(memSize);
cudaMalloc( (void**) &d_a, memSize);
```

// launch kernel

```
dim3 dimGrid(num_blocks);
dim3 dimBlock(num_th_per_blk);
initArray<<< dimGrid, dimBlock >>>(d_a);
```

// retrieve results

```
cudaMemcpy( h_a, d_a, memSize,
            cudaMemcpyDeviceToHost);
```

...

// kernel – find linearized threadId, and set A[tid] = tid

```
__global__ void initArray( int *A)
{
    int tid;

    tid      = blockIdx.x * blockDim.x + threadIdx.x;
    A[ tid ] = tid;
}
```

initializes an array according
to the linearized thread_id

← for this example, assume we will
have one thread per data item

it's your turn . . . reverseArray_singleBlock

// initialize an array using GPU

```
int main()
{
    int *h_a;           //pointer for host memory
    int *d_a;           //pointer for device memory
```

// define thread hierarchy

```
int num_blocks      = 8;
int num_th_per_blk  = 8;
```

// allocate host and device memory

```
size_t memSize;
```

```
memSize = num_blocks * num_th_per_blk * sizeof(int);
h_a      = (int*) malloc(memSize);
cudaMalloc( (void**) &d_a, memSize);
```

// launch kernel

```
dim3 dimGrid(num_blocks);
dim3 dimBlock(num_th_per_blk);
initArray<<< dimGrid, dimBlock >>>(d_a);
```

// retrieve results

```
cudaMemcpy( h_a, d_a, memSize,
            cudaMemcpyDevice To Host);
```

...

// kernel – find linearized threadIdx, and set A[tid] = tid

```
__global__ void initArray( int *A)
{
    int tid;

    tid      = blockIdx.x * blockDim.x + threadIdx.x;
    A[ tid ] = tid;
}
```

modify host and initArray kernel to:

- use a single block of 256 threads
- use an array of 256 integers
- on host, initialize the values of an array to equal the array offset
- in the kernel, reverse the contents of the array
- on the host, verify the kernel's result

reverseArray_singleBlock

IYT answer

cse5441 - parallel computing

https://cdn8.bigcommerce.com/s-o1bh36yff/images/stencil/1280x1280/products/18124/29821/hpgraphiccardsvideocardtesla__52934.1516995900.jpg?c=2&imbypass=on
<https://i0.wp.com/www.visitcocosisland.com/wp-content/uploads/2017/01/animalia-life-club.jpg?resize=640%2C492>
<https://qph.fs.quoracdn.net/main-qimg-257ea3b2d296c1605a50e5fd5db21cfa-c>
<https://pictures.topspeed.com/IMG/jpg/201205/plymouth-hemi-cuda-p.jpg>



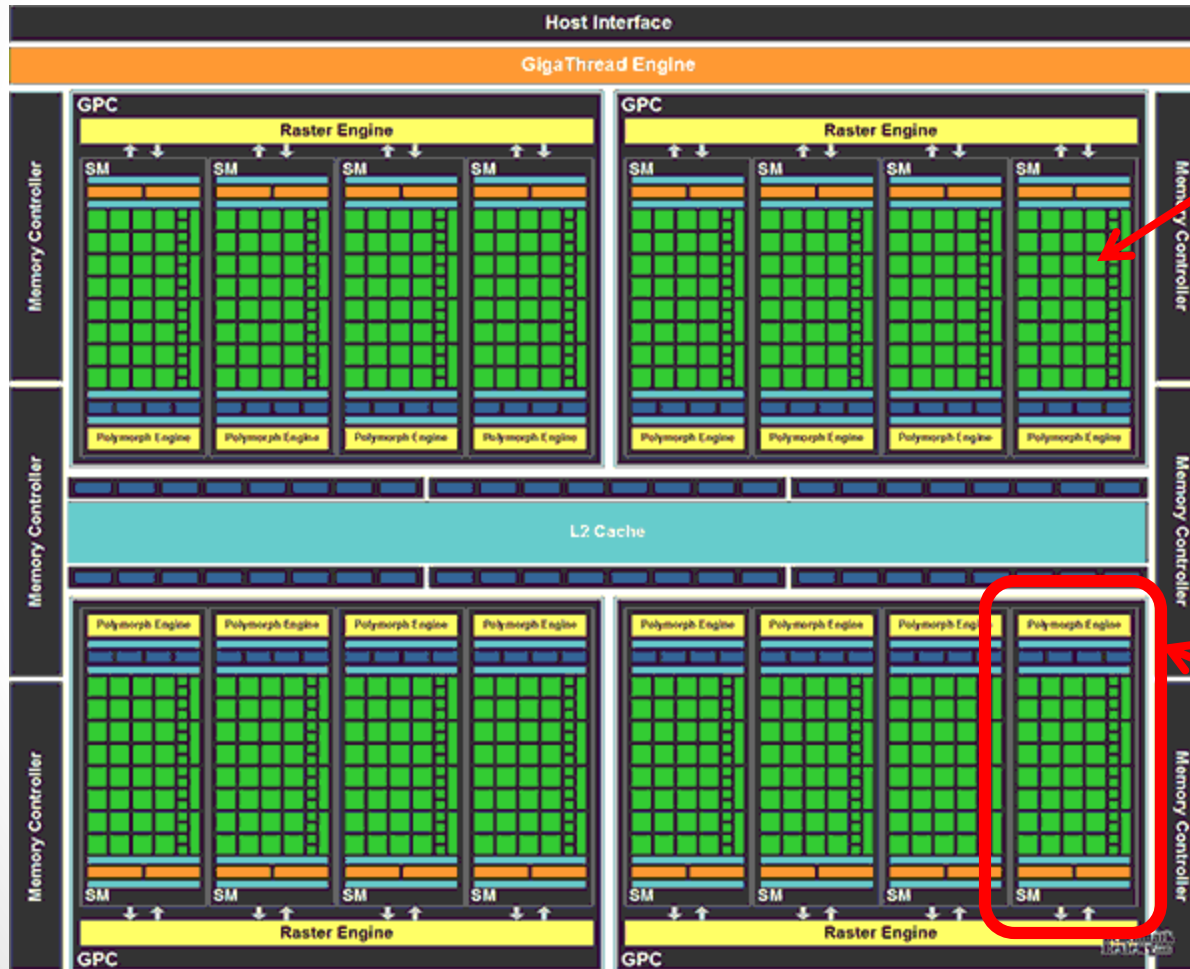
‘cuda



OSU CSE 2321

$$\begin{bmatrix} 25 \end{bmatrix}$$

“massively parallel” thread organization: threads



a thread runs
on a core

cores are
components of
streaming
multiprocessors
(SM)

“massively parallel” thread organization: blocks

currently all blocks from the same grid run on the same SM
plans are for block scheduling across SMs



threads are
organized into
blocks

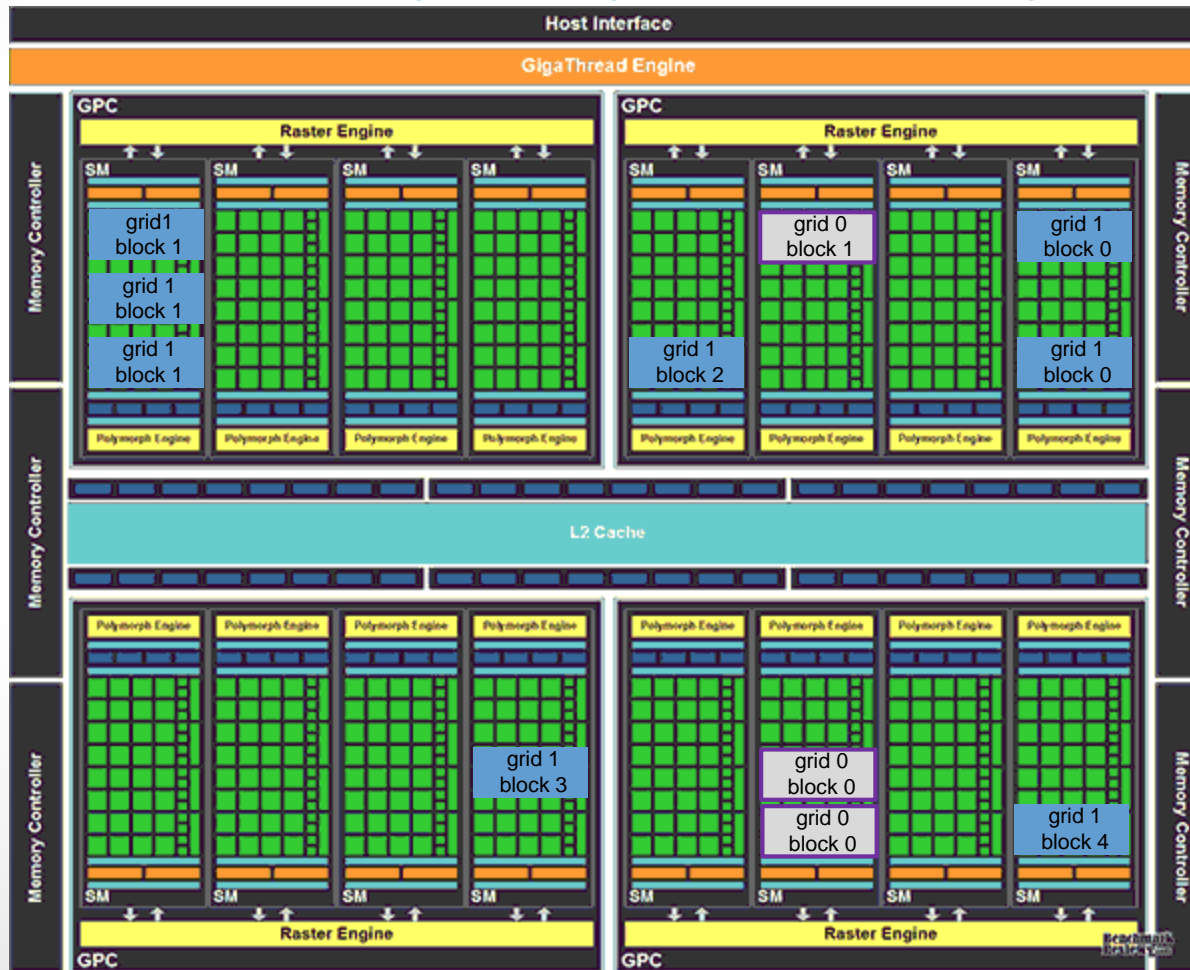
threads within
a **block** are
organized into
warps
of 32 threads

warps are
scheduled
atomically
and all **block**
warps run on the
same **streaming**
multiprocessor (SM)

note: warps are stylized for illustration purposes

“massively parallel” thread organization: grids

warps from 2 different grids running simultaneously on a single GPU




a group of **blocks**
is a **grid**

each **kernel**
gets one **grid**
of **threads**

a **GPU** may
schedule
multiple
simultaneous
grids

“massively parallel” thread organization



“massively parallel” thread organization:

threads in blocks, 1 block

1D: threadIdx.x

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

2D: threadIdx.x * threadIdx.y

0,0	0,1	0,2	0,3	0,4	0,5	0,6	0,7
1,0	1,1	1,2	1,3	1,4	1,5	1,6	1,7
2,0	2,1	2,2	2,3	2,4	2,5	2,6	2,7

3D: threadIdx.x * threadIdx.y * threadIdx.z

0,0,0	0,0,1	0,0,2	0,0,3	0,0,4	0,0,5	0,0,6	0,0,7			
0,1,0	0,1,1	0,1,2	0,1,3	0,1,4	0,1,5	0,1,6	0,1,7	,7		
0,2,0	0,2,1	0,2,2	0,2,3	0,2,4	0,2,5	0,2,6	0,2,7	,7	,7	
1,2,0	1,2,1	1,2,2	1,2,3	1,2,4	1,2,5	1,2,6	1,2,7	,7	,7	,7
2,2,0	2,2,1	2,2,2	2,2,3	2,2,4	2,2,5	2,2,6	2,2,7			

threadIdx.x

thread position in dimension x

threadIdx.y

thread position in dimension y

threadIdx.z

thread position in dimension z

1-dimensional:

```
dim3 dimBlock(n_th_per_blk);
blk_t_id = threadIdx.x
```

2-dimensional:

```
dim3 dimBlock(nth_x, nth_y);
blk_t_id = blockDim.x * threadIdx.y
          + threadIdx.x
```

3-dimensional:

```
dim3 dimBlock(nth_x, nth_y, nth_z);
blk_t_id = threadIdx.z * blockDim.y * blockDim.x
          + blockDim.x * threadIdx.y
          + threadIdx.x
```

“massively parallel” thread organization:

threads in blocks, 1D x 1D

1D: blockIdx.x = 0

1D: threadIdx.x

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

1D: blockIdx.x = 1

1D: threadIdx.x

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

1D: blockIdx.x = 2

1D: threadIdx.x

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

blockIdx.x

block position in dimension x

blockIdx.y

block position in dimension y

blockIdx.z

block position in dimension z

1-dimensional x 1-dimensional:


```
dim3 dimBlock(n_th_per_blk);  
dim3 dimGrid(n_blk_per_grid);
```

```
blockNumInGrid   = blockIdx.x  
threadsPerBlock  = blockDim.x  
threadNumInBlock = threadIdx.x
```

```
global_t_id = blockIdx.x * blockDim.x  
            + threadIdx.x
```

“massively parallel” thread organization:

threads in blocks, 2D x 2D



`blockIdx.x`

block position in dimension x

`blockIdx.y`

block position in dimension y

// if we had a third dimension

(`blockIdx.z` block position in dimension z)

2-dimensional x 2-dimensional:

`dim3 dimBlock(tpb_x, tpb_y);`

`dim3 dimGrid(bpg_x, bpg_y);`

`blockNumInGrid = blockIdx.x +`
`gridDim.x * blockIdx.y`

`threadsPerBlock = blockDim.x * blockDim.y`

`threadNumInBlock = threadIdx.x +`
`blockDim.x * threadIdx.y`

`global_t_id = blockNumInGrid`
`* threadsPerBlock`
`+ threadNumInBlock`

note: this illustration in column-major

it's your turn: reverseArray_multiBlock1

```
// reverse an array using GPU
```

```
int main()
{
    int *h_a;           //pointer for host memory
    int *h_v;           //pointer verification copy
    int *d_a;           //pointer for device input
    int *d_b;           //pointer for device output
```

```
// define thread hierarchy
```

```
int nblocks   = 1;
int dimA      = 256;
int tpb       = dimA;
```

```
// allocate host and device memory
```

```
size_t memSize;
memSize = nblocks * tpb * sizeof(int);
h_a      = (int*) malloc(memSize);
h_v      = (int*) malloc(memSize);
cudaMalloc( (void**) &d_a, memSize);
cudaMalloc( (void**) &d_b, memSize);
```

```
// initialize host arrays, copy to device
```

```
for (int i = 0; i < dimA; i++)
```

```
{
    h_a[i] = i;
    h_v[i] = i;
}
```

```
cudaMemcpy(d_a, h_a, memSize,
            cudaMemcpyHostToDevice
```

```
// launch kernel
```

```
dim3 dimGrid(nblocks);
dim3 dimBlock(tpb);
reverseArrayMBlock1<<< dimGrid, dimBlock >>>(d_a, d_b);
```

```
// get results
```

```
cudaMemcpy( h_a, d_b, memSize,
             cudaMemcpyDevice To Host);
```

```
...
```

```
// kernel – reverse inArray into outArray
```

```
__global__ void reverseArrayMBlock1( int *inArray, int *outArray)
{
    int tid;
```

```
    outArray[blockDim.x – 1 – threadIdx.x] = inArray[threadIdx.x]
```

modify host and kernel to:

- use 2 blocks of 1024 threads each
- use a linear array of 1M (1024x1024) integers
- on host, initialize the values of the array to equal the array offset
- in the kernel, reverse the contents of the array
- on the host, verify the kernel's result

reverseArrayMultiBlock1

IYT answer

it's your turn: reverseArray_Block2

```
// reverse an array using GPU
```

```
int main()
{
    int *h_a;           //pointer for host memory
    int *h_v;           //pointer verification copy
    int *d_a;           //pointer for device input
    int *d_b;           //pointer for device output
    int dimA = 1024*1024; //size of array
```

```
// define thread hierarchy
```

```
int nblocks = 2;
int tpb = 1024;
```

```
// allocate host and device memory
```

```
size_t memSize;
memSize = nblocks * tpb * sizeof(int);
h_a = (int*) malloc(memSize);
h_v = (int*) malloc(memSize);
cudaMalloc( (void**) &d_a, memSize);
cudaMalloc( (void**) &d_b, memSize);
```

```
// initialize host arrays, copy to device
```

```
for (int i = 0; i < dimA; i++)
```

```
{
    h_a[i] = i;
    h_v[i] = i;
}
```

```
cudaMemcpy(d_a, h_a, memSize,
           cudaMemcpyHostToDevice)
```

```
// launch kernel
```

```
dim3 dimGrid(nblocks);
dim3 dimBlock(tpb);
reverseArrayBlock2<<< dimGrid, dimBlock >>>(d_a, d_b);
```

```
// get results
```

```
cudaMemcpy( h_a, d_b, memSize,
            cudaMemcpyDevice To Host);
```

```
...
```

```
// kernel – reverse inArray into outArray
```

```
__global__ void reverseArray_Block2( int *inArray, int *outArray)
{
    int tid;
```

```
    outArray[blockDim.x – 1 – threadIdx.x] = inArray[threadIdx.x]
```

in the kernel, reverse the contents of the array
as before, except:

- use a single block of 1024 threads
- use only a single on-device array

reverseArrayMultiBlock2

IYT answer

matrix multiply example

CPU implementation

```
int main()
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    MatrixMulHost( A, B, C, dim);
}

void MatrixMulHost (float *A, float *B, float *C, int dim)
{
    float a, b, sum;

    for (int i = 0; i < dim; i++)
    {
        for (int j = 0; j < dim; j++)
        {
            sum = 0;
            for (int k = 0; k < dim; k++)
            {
                a    = A[ i * dim + k ];
                b    = B[ k * dim + j ];
                sum += a * b;
            }
            C[ i * dim + j ] = sum;
        }
    }
}
```

matrix multiply example

GPU skeleton

```
int main()
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B ...
    MatrixMulHost( A, B, C, dim);
}

void MatrixMulHost (float *A, float *B, float *C, int dim)
{
    float a, b, sum;

    for (int i = 0; i < dim; i++)
    {
        for (int j = 0; j < dim; j++)
        {
            sum = 0;
            for (int k = 0; k < dim; k++)
            {
                a    = A[ i * dim + k ];
                b    = B[ k * dim + j ];
                sum += a * b;
            }
            C[ i * dim + j ] = sum;
        }
    }
}
```

```
int main(void)
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B ...

    // perform MatrixMul on Device

}
```

matrix multiply example

GPU skeleton

```
int main()
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    MatrixMulHost( A, B, C, dim);
}

void MatrixMulHost (float *A, float *B, float *C, int dim)
{
    float a, b, sum;

    for (int i = 0; i < dim; i++)
    {
        for (int j = 0; j < dim; j++)
        {
            sum = 0;
            for (int k = 0; k < dim; k++)
            {
                a    = A[ i * dim + k ];
                b    = B[ k * dim + j ];
                sum += a * b;
            }
            C[ i * dim + j ] = sum;
        }
    }
}
```

```
int main(void)
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .

    // define thread hierarchy
    // allocate device memory
    // initialize device memory
    // launch kernel
    // perform MatrixMul on Device
    // retrieve results

}
```

matrix multiply example

GPU thread hierarchy

```
int main()
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    MatrixMulHost( A, B, C, dim);
}

void MatrixMulHost (float *A, float *B, float *C, int dim)
{
    float a, b, sum;

    for (int i = 0; i < dim; i++)
    {
        for (int j = 0; j < dim; j++)
        {
            sum = 0;
            for (int k = 0; k < dim; k++)
            {
                a    = A[ i * dim + k ];
                b    = B[ k * dim + j ];
                sum += a * b;
            }
            C[ i * dim + j ] = sum;
        }
    }
}
```

```
int main(void)
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .

    // define thread hierarchy
    int nblocks    = 4;
    int tpb        = 512;

    // allocate device memory
    // initialize device memory
    // launch kernel
    // perform MatrixMul on Device
    // retrieve results
}
```

matrix multiply example

GPU memory allocation

```
int main()
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    MatrixMulHost( A, B, C, dim);
}

void MatrixMulHost (float *A, float *B, float *C, int dim)
{
    float a, b, sum;

    for (int i = 0; i < dim; i++)
    {
        for (int j = 0; j < dim; j++)
        {
            sum = 0;
            for (int k = 0; k < dim; k++)
            {
                a    = A[ i * dim + k ];
                b    = B[ k * dim + j ];
                sum += a * b;
            }
            C[ i * dim + j ] = sum;
        }
    }
}
```

```
int main(void)
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .

    // define thread hierarchy
    int nblocks    = 4;
    int tpb        = 512;

    // allocate device memory
    size_t memSize;
    memSize = dim * dim * sizeof(float);
    cudaMalloc( (void**) &d_A, memSize);
    cudaMalloc( (void**) &d_B, memSize);
    cudaMalloc( (void**) &d_C, memSize);

    // initialize device memory
    // launch kernel
    // perform MatrixMul on Device
    // retrieve results

}
```

matrix multiply example

GPU memory initialization

```
int main()
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    MatrixMulHost( A, B, C, dim);
}

void MatrixMulHost (float *A, float *B, float *C, int dim)
{
    float a, b, sum;

    for (int i = 0; i < dim; i++)
    {
        for (int j = 0; j < dim; j++)
        {
            sum = 0;
            for (int k = 0; k < dim; k++)
            {
                a    = A[ i * dim + k ];
                b    = B[ k * dim + j ];
                sum += a * b;
            }
            C[ i * dim + j ] = sum;
        }
    }
}
```

```
int main(void)
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    // define thread hierarchy
    int nblocks    = 4;
    int tpb        = 512;
    // allocate device memory
    size_t memSize;
    memSize = dim * dim * sizeof(float);
    cudaMalloc( (void**) &d_A, memSize);
    cudaMalloc( (void**) &d_B, memSize);
    cudaMalloc( (void**) &d_C, memSize);

    // initialize device memory
    cudaMemcpy(d_A, A, memSize, cudaMemcpy HostToDevice);
    cudaMemcpy(d_B, B, memSize, cudaMemcpy HostToDevice);

    // launch kernel
    // perform MatrixMul on Device
    // retrieve results

}
```

matrix multiply example

GPU kernel launch

```
int main()
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    MatrixMulHost( A, B, C, dim);
}

void MatrixMulHost (float *A, float *B, float *C, int dim)
{
    float a, b, sum;

    for (int i = 0; i < dim; i++)
    {
        for (int j = 0; j < dim; j++)
        {
            sum = 0;
            for (int k = 0; k < dim; k++)
            {
                a      = A[ i * dim + k ];
                b      = B[ k * dim + j ];
                sum += a * b;
            }
            C[ i * dim + j ] = sum;
        }
    }
}
```

```
int main(void)
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    // define thread hierarchy
    int  nblocks   = 4;
    int  tpb       = 512;
    // allocate device memory
    size_t memSize;
    memSize = dim * dim * sizeof(float);
    cudaMalloc( (void**) &d_A, memSize);
    cudaMalloc( (void**) &d_B, memSize);
    cudaMalloc( (void**) &d_C, memSize);
    // initialize device memory
    cudaMemcpy(d_A, A, memSize, cudaMemcpy HostToDevice);
    cudaMemcpy(d_B, B, memSize, cudaMemcpy HostToDevice);

    // launch kernel
    dim3 dimGrid(nblocks);
    dim3 dimBlock(tpb);
    MatrixMulDevice<<< dimGrid, dimBlock >>>(d_A, d_B, d_C, dim);
    // perform MatrixMul on Device

    // retrieve results

}
```

matrix multiply example

fetch results

```
int main()
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    MatrixMulHost( A, B, C, dim);
}

void MatrixMulHost (float *A, float *B, float *C, int dim)
{
    float a, b, sum;

    for (int i = 0; i < dim; i++)
    {
        for (int j = 0; j < dim; j++)
        {
            sum = 0;
            for (int k = 0; k < dim; k++)
            {
                a      = A[ i * dim + k ];
                b      = B[ k * dim + j ];
                sum += a * b;
            }
            C[ i * dim + j ] = sum;
        }
    }
}
```

```
int main(void)
{
    float *A, *B, *C; int dim = 512;
    A = new float(dim*dim); B = new float(dim*dim);
    C = new float(dim*dim);
    // I/O to load A and B .
    // define thread hierarchy
    int  nblocks    = 4;
    int  tpb        = 512;
    // allocate device memory
    size_t memSize;
    memSize = dim * dim * sizeof(float);
    cudaMalloc( (void**) &d_A, memSize);
    cudaMalloc( (void**) &d_B, memSize);
    cudaMalloc( (void**) &d_C, memSize);
    // initialize device memory
    cudaMemcpy(d_A, A, memSize, cudaMemcpy HostToDevice);
    cudaMemcpy(d_B, B, memSize, cudaMemcpy HostToDevice);

    // launch kernel
    dim3 dimGrid(nblocks);
    dim3 dimBlock(tpb);
    MatrixMulDevice<<< dimGrid, dimBlock >>>(d_A, d_B, d_C, dim);
    // perform MatrixMul on Device

    // retrieve results
    cudaMemcpy( C, d_C, memSize, cudaMemcpyDevice To Host);
}
```


it's your turn: matrix multiply example

kernel

```
void MatrixMulHost (float *A, float *B, float *C, int dim)
{
float a, b, sum;

for (int i = 0; i < dim; i++)
{
for (int j = 0; j < dim; j++)
{
sum = 0;
for (int k = 0; k < dim; k++)
{
a = A[ i * dim + k ];
b = B[ k * dim + j ];
sum += a * b;
}
C[ i * dim + j ] = sum;
}
}
}
```

```
__global__ void MatrixMulDevice( float *A, float *B, float *C)
{
// perform MatrixMul on Device
```

provide the kernel

```
int main(void)
{
float *A, *B, *C; int dim = 512;
A = new float(dim*dim); B = new float(dim*dim);
C = new float(dim*dim);
// I/O to load A and B .
// define thread hierarchy
int nblocks = 4;
int tpb = 512;
// allocate device memory
size_t memSize;
memSize = dim * dim * sizeof(float);
cudaMalloc( (void**) &d_A, memSize);
cudaMalloc( (void**) &d_B, memSize);
cudaMalloc( (void**) &d_C, memSize);
// initialize device memory
cudaMemcpy(d_A, A, memSize, cudaMemcpy HostToDevice);
cudaMemcpy(d_B, B, memSize, cudaMemcpy HostToDevice);

// launch kernel
dim3 dimGrid(nblocks);
dim3 dimBlock(tpb);
// perform MatrixMul on Device
MatrixMulDevice<<< dimGrid, dimBlock >>>(d_A, d_B, d_C, dim);

// retrieve results
cudaMemcpy( C, d_C, memSize, cudaMemcpyDevice To Host);
}
```

matrix multiply example

IYT answer

engineering cuda kernels

limits vary according to Nvidia "compute capability" level

- do not exceed thread/block limit (512 or 1024)
- do not exceed thread/block dimensionality (512,512,64 or 1024,1024,64)
- do not exceed total registers (8k, 16k, 32k or 64k)
- do not exceed block shared memory (16kB or 48kB)

top **performance** determined by benchmarking

- depends upon specific GPU hardware
- depends upon application characteristics
- threads execute in 32-thread warps
- typical sweet spot in the 128 – 512 threads/block range


thread “divergence”

```
// kernel – do various things to array A

__global__ void doLotsOfStuff( float *A)
{
    int tid;

    tid = blockIdx.x * blockDim.x + threadIdx.x;
    if ( tid % 2 == 0 )
        A[tid] = 42.0;
    if ( tid % 3 == 0 )
        A[tid] = 2.0;
    if ( threadIdx.x > ( th_per_warp / 2 ) )
        A[tid] = pow(A[tid], 3);
        A[tid] = A[tid] / A[tid-1] + A[tid-2];
    A[tid] - - ;
    if ( A[tid] <= 2 )
        A[tid] = abs(A[tid]) + 1;
    if ( tid % 2 == 0 )
        A[tid] = A[tid] + 12.0;
        A[tid] = A[tid] * PI;
    else
        A[tid] = abs(A[tid] - 12.0);
        A[tid] = A[tid] / PI;
}
```

active threads in warp



in-active threads in warp

SIMD serialization

arbitrarily-ordered data

blue	heavy
blue	light
green	light
red	heavy
blue	medium
blue	light
green	light
red	heavy
blue	heavy
blue	heavy
green	medium
red	medium
blue	light
red	light
green	light
green	medium
red	medium
...	

naturally lends itself to a branching structure ...

```
if blue
    if heavy
        do blue-heavy work
    else if medium
        do blue-medium work
    else if light
        do blue-light work
else if green
    if heavy
        do green-heavy work
    else if medium
        do green-medium work
    else if light
        do green-light work
else if red
    ...
```

SIMD serialization

SIMD-ordered data

blue	heavy
blue	heavy
blue	heavy
...	
blue	medium
blue	medium
blue	medium
...	
blue	light
blue	light
blue	light
...	
green	heavy
green	heavy
green	heavy
...	
green	medium
...	

naturally lends itself to a branching structure ...

```
while blue heavy
    do blue-heavy work
while blue medium
    do blue-medium work
while blue light
    do blue-light work
while green heavy
    do green-heavy work
while green medium
    do green-medium work
while green light
    do green-light work
while red heavy
    do red-heavy work
while red medium
    ...
```

giga flops (GFlops) benchmarking

given that all variables are single- or double-precision floating point numbers,
the following statements perform "1 flop" worth of work:

- $A = B + C$
- $A[k] = A[j] * A[i]$
- $A[k-1] += B[m]$
- $A[k] = B[m] / C[n]$

while the following would be "3 flops" worth of work:

- $A[m] = (B[m] + C[m]) / (D[m] + E[m])$

to compute the performance of a code segment:

- count the number of flops performed in the segment
- divide by the execution time in seconds

results are typically reported in "GFlops/sec"

cuda function types

`__global__ void`

kernel callable from non-cuda function

example:

```
__global__ void CellLoop1(struct cuda_params *p, float *anbs, float *aps, float *resid, float *scs)
```

`__host__ void`

function (non-kernel) using cuda functions
callable from non-cuda function

example:

```
__host__ void disord_simd(int tindex, int oiter)
//useful for including operations like ...
    cudaMemcpy(anbs, d_anbs, ncells*nf_max*sizeof(float), cudaMemcpyDeviceToHost);
```

`__device__ float`

function callable from a kernel

example:

```
__device__ float d_iwalls(int pol, int band, float degrees, float delta_la, int nla, float *xi, float *wi);
```

cuda functions reside in a “.cu” program file

debugging: gpuAssert()

host-side debugging

a helpful macro ...

```
#define gpuErrchk(ans) { gpuAssert((ans), __FILE__, __LINE__); }
inline void gpuAssert(cudaError_t code, const char *file, int line, bool abort=true)
{
    if (code != cudaSuccess)
    {
        fprintf(stderr,"GPUassert: %s %s %d\n", cudaGetErrorString(code), file, line);
        if (abort) exit(code);
    }
}
```

verify copy successful

```
cudaMemcpy(r, d_resid, ncells*sizeof(float), cudaMemcpyDeviceToHost);
gpuErrchk( cudaMemcpy(r, d_resid, ncells*sizeof(float), cudaMemcpyDeviceToHost) );
```

ensure kernel ran to completion

```
CellLoop3<<<dimGrid, dimBlock>>>(d_p, d_anbs, d_aps, d_intensity, d_resid, d_scs);
gpuErrchk( cudaPeekAtLastError() );           // stop if kernel did not launch
gpuErrchk( cudaDeviceSynchronize() );         // for debugging – ensure kernel finished
```

debugging: printf()

device-side debugging

“printf” from within a kernel is supported on selected cuda implementations

- must have compute capability of 2.0 or higher
- must have cuda developer version 3.1 or higher
- output is buffered and returned to host at kernel completion
- output buffer overflow is lost
- host program must run long enough for output buffers to transfer and flush (think seg fault behavior)

debugging

- cudagrind
- cuda-gdb

cse5441 - parallel computing

https://cdn8.bigcommerce.com/s-o1bh36yff/images/stencil/1280x1280/products/18124/29821/hpgraphiccardsvideocardtesla__52934.1516995900.jpg?c=2&imbypass=on
<https://i0.wp.com/www.visitcocosisland.com/wp-content/uploads/2017/01/animalia-life-club.jpg?resize=640%2C492>
<https://qph.fs.quoracdn.net/main-qimg-257ea3b2d296c1605a50e5fd5db21cfa-c>
<https://pictures.topspeed.com/IMG/jpg/201205/plymouth-hemi-cuda-p.jpg>



‘cuda

