

cse5441 - parallel computing

advanced instrumentation

Performance API (PAPI)

- discussion standard for hardware-enabled performance measurement
- cross-platform support
- potential for portable solution
- incentive for multi-vendor support
- ease-of-use



PAPI interfaces

high-level interface

- start, read and stop event counters

<u>Function name</u>	<u>Description</u>
PAPI_num_counters()	get the number of hardware counters available on the system
PAPI_flips ()	Mflips/s (floating point instruction rate)
PAPI_flops()	Mflops/s (floating point operation rate)
PAPI_ipc()	gets instructions per cycle
PAPI_accum_counters()	add current counts to array and reset counters
PAPI_read_counters()	copy current counts to array and reset counters
PAPI_start_counters()	start counting hardware events
PAPI_stop_counters()	stop counters and return current counts

PAPI standard events

examples

PAPI_L1_DCM	Level 1 data cache misses
PAPI_L1_ICM	Level 1 instruction cache misses
PAPI_L2_DCM	Level 2 data cache misses
PAPI_L2_ICM	Level 2 instruction cache misses
PAPI_L3_DCM	Level 3 data cache misses
PAPI_L3_ICM	Level 3 instruction cache misses
PAPI_L1_TCM	Level 1 total cache misses
PAPI_L2_TCM	Level 2 total cache misses
PAPI_L3_TCM	Level 3 total cache misses
...	
PAPI_LD_INS	Load instructions executed
PAPI_SR_INS	Store instructions executed
PAPI_L1_LDM	Level 1 load misses
PAPI_L1_STM	Level 1 store misses
PAPI_L2_LDM	Level 2 load misses
PAPI_L2_STM	Level 2 store misses
...	
PAPI_TOT_CYC	Total cycles
PAPI_TOT_INS	Total instructions executed
PAPI_FP_INS	Floating point instructions executed
...	

PAPI cache metrics

L1 data cache hit rate:

$$1.0 - (\text{PAPI_L1_DCM} / (\text{PAPI_LD_INS} + \text{PAPI_SR_INS}))$$

L2 data cache hit rate:

$$1.0 - (\text{PAPI_L2_DCM} / \text{PAPI_L1_DCM})$$

basic PAPI demo

Function name

PAPI_start_counters()

PAPI_stop_counters()

Description

start counting hardware events

stop counters and return current counts

```

/*****\
 * Initialize the PAPI library and get the number of counters available *
\*****/
int      Events[2] = { PAPI_TOT_CYC, PAPI_TOT_INS };
int      num_hwcntrs = 0;
long_long values[2];

if ((num_hwcntrs = PAPI_num_counters()) <= PAPI_OK)
{
    cerr << "PAPI init error -- exiting" << endl;
    exit(1);
}
cout << "This system has " << num_hwcntrs << " available counters.\n";

t = clock();
if (PAPI_start_counters(Events, 2) != PAPI_OK)
{
    cerr << "PAPI start error -- exiting" << endl;
    exit(1);
}
/*****/
 * run a simple stencil computation on the grid *
\*****/
Mesh.dissipate(Mesh.affect_rate, Mesh.epsilon);
    
```

```

/*****/
 * print results *
\*****/
if (PAPI_stop_counters(values, 2) != PAPI_OK)
{
    cerr << "PAPI count error -- exiting" << endl;
    exit(1);
}

t = clock() - t;
cout << "Convergence time reported by clock(): " << setw(20) << t << endl;
cout << "          PAPI total cycles: " << setw(20) << values[0] << endl;
cout << "          PAPI total instructions: " << setw(20) << values[1] << endl;
    
```

PAPI interfaces

low-level interface

- customizable access to event counters

```
#include <papi.h>
#include <stdio.h>
```

```
main()
{
    int retval, EventSet = PAPI_NULL;
    long_long values[1];
```

```
/* Initialize the PAPI library */
retval = PAPI_library_init(PAPI_VER_CURRENT);
```

```
if (retval != PAPI_VER_CURRENT) {
    fprintf(stderr, "PAPI library init error!\n");
    exit(1);
}
```

```
/* Create the Event Set */
if (PAPI_create_eventset(&EventSet) != PAPI_OK)
    handle_error(1);
```

```
/* Add Total Instructions Executed to our EventSet */
if (PAPI_add_event(EventSet, PAPI_TOT_INS) != PAPI_OK)
    handle_error(1);
```

```
/* Start counting */
if (PAPI_start(EventSet) != PAPI_OK)
    handle_error(1);

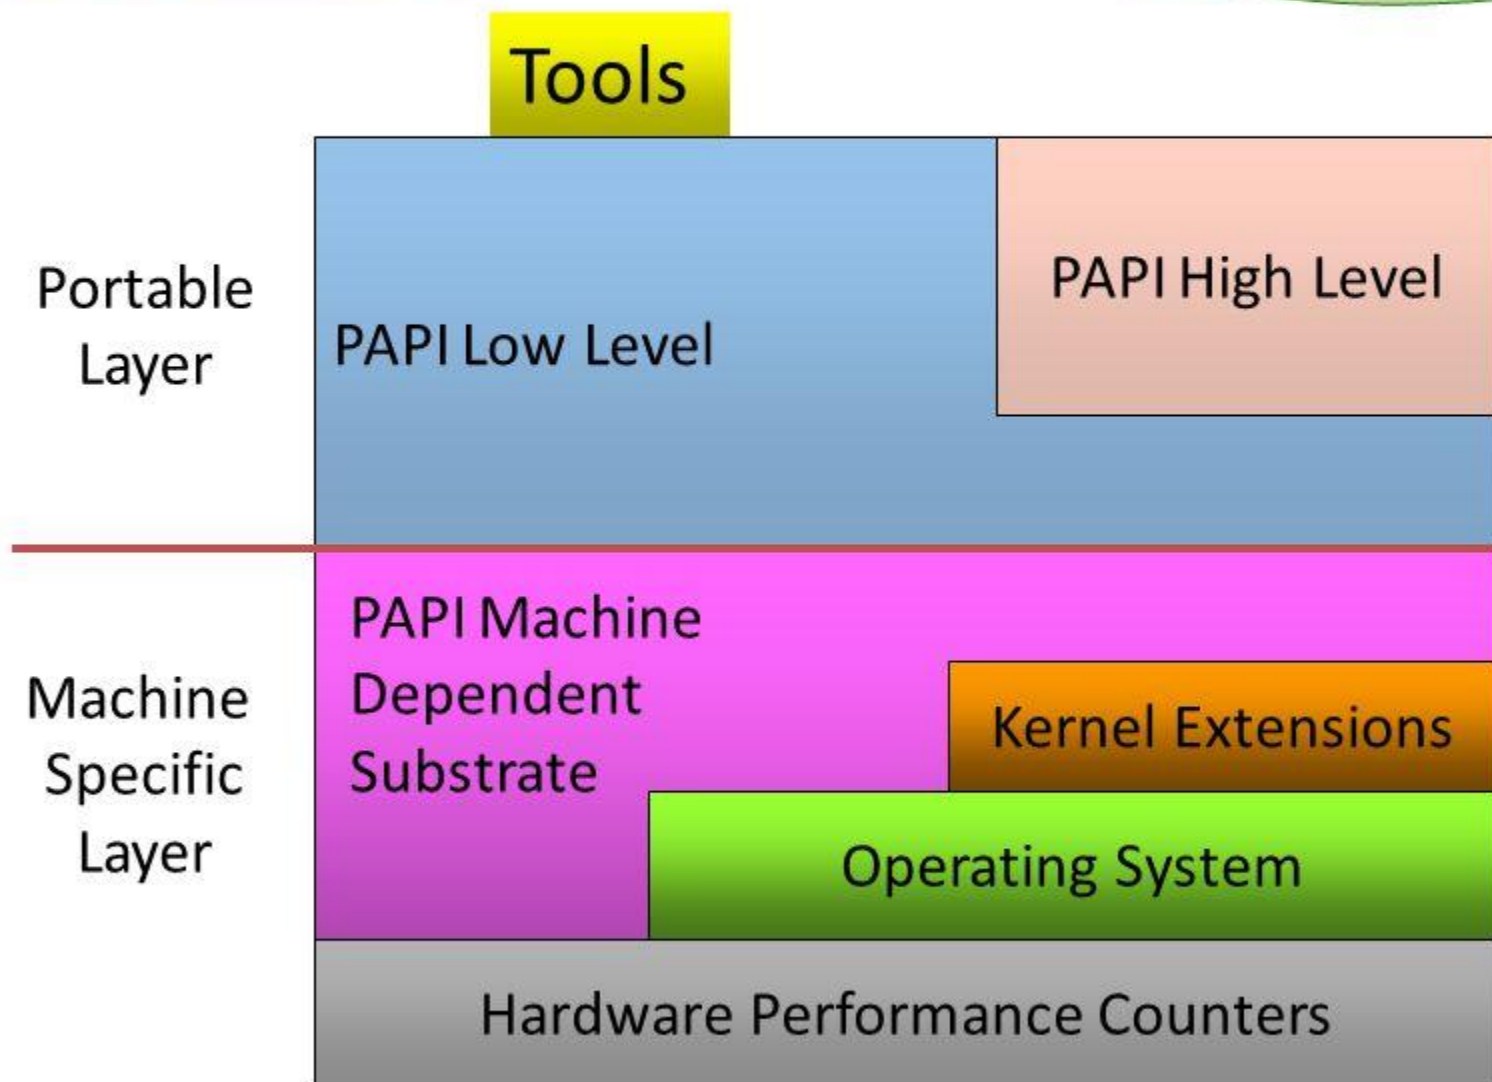
/* Do some computation here */

if (PAPI_read(EventSet, values) != PAPI_OK)
    handle_error(1);

/* Do some computation here */

if (PAPI_stop(EventSet, values) != PAPI_OK)
    handle_error(1);
}
```

PAPI Architecture



PAPI thread support

- must use kernel (or “bound”) threads for **thread-level** measurement
 - PTHREAD_SCOPE_SYSTEM for Pthreads
 - otherwise, all threads report process values

```
include <papi.h>
#include <pthread.h>
```

```
main()
{
    unsigned long int tid;
```

```
    if (PAPI_library_init(PAPI_VER_CURRENT) != PAPI_VER_CURRENT)
        exit(1);
```

← uses low-level interface

```
    if (PAPI_thread_init(pthread_self) != PAPI_OK)
        exit(1);
```

← each thread must initialize

```
    if ((tid = PAPI_thread_id()) == (unsigned long int)-1)
        exit(1);
```

← can now request thread id

```
    printf("Initial thread id is: %lu\n",tid);
}
```

← will be 0 for master thread

PAPI thread support

<code>PAPI_register_thread()</code>	each thread must register
<code>PAPI_unregister_thread()</code>	release prior to reuse
<code>PAPI_get_thr_specific("tag, ptr")</code>	four defined memory locations for thread-specific data (why 4?)
<code>PAPI_set_thr_specific("tag, ptr")</code>	

PAPI portability

- goal is not to compare systems
- implementation of as many standard events as possible
 - while avoiding misleading or erroneous results
- vendor counter interfaces sometimes have bugs
 - counter overflow

PAPI on OSC

- version 5.4.1 available on [Oakley](#) and [Ruby](#)
 - module load gnu/4.8.5
 - module load papi
- `g++ -lpapi`

PAPI cautions

- not all functions in both C and Fortran
- documentation is inconsistent
- not all functions supported on all platforms
- some errors in hardware counters
- no commercial support

PAPI further reading

Carmen:

- | | |
|---------------------------|------------------------|
| • papi-journal-final_2000 | original journal paper |
| • papi-conference-2001 | similar, more brief |
| • papi-dongarra-2003 | update |

On-line documentation:

- http://icl.cs.utk.edu/projects/papi/wiki/Main_Page

cse5441 - parallel computing

advanced instrumentation